

```
!pip install --upgrade scikit-learn

Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (1.2.2)
Collecting scikit-learn
  Downloading scikit_learn-1.4.1.post1-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (12.1 MB)
    12.1/12.1 MB 53.9 MB/s eta 0:00:00
Requirement already satisfied: numpy<2.0,>=1.19.5 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.25.2)
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.11.4)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (3.4.0)
Installing collected packages: scikit-learn
  Attempting uninstall: scikit-learn
    Found existing installation: scikit-learn 1.2.2
    Uninstalling scikit-learn-1.2.2:
      Successfully uninstalled scikit-learn-1.2.2
Successfully installed scikit-learn-1.4.1.post1
```

```
import sklearn
print(sklearn.__version__)
```

1.4.1.post1

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

```
#Load the data
from google.colab import files
uploaded = files.upload()
```

Choose Files

No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving IBM HR Final cleaned Data.xlsm to IBM HR Final cleaned Data.xlsm

```
df=pd.read_excel('IBM_HR_Final_cleaned_Data.xlsm')
```

```
df.head()
```

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Environm
0	49	Travel_Frequently	279	Research & Development	8	1	
1	59	Non-Travel	1420	Human Resources	2	4	
2	59	Non-Travel	1420	Human Resources	2	4	
3	49	Travel_Frequently	279	Research & Development	8	1	
4	49	Travel_Frequently	279	Research & Development	8	1	
5 rows × 30 columns							

```
df.tail()
```

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Envi
19473	47	Travel_Rarely	465	Research & Development	1	3	
19474	38	Travel_Rarely	371	Research & Development	2	3	
19475	34	Travel_Rarely	629	Research & Development	27	2	
19476	55	Non-Travel	177	Research & Development	8	1	
19477	27	Travel_Rarely	1134	Research & Development	16	4	

5 rows × 30 columns

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19478 entries, 0 to 19477
Data columns (total 30 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    19478 non-null  int64
1   BusinessTravel                        19478 non-null  object
2   DailyRate                            19478 non-null  int64
3   Department                            19478 non-null  object
4   DistanceFromHome                     19478 non-null  int64
5   Education                            19478 non-null  int64
6   EnvironmentSatisfaction               19478 non-null  int64
7   Gender                                19478 non-null  object
8   HourlyRate                           19478 non-null  int64
9   JobInvolvement                       19478 non-null  int64
10  JobLevel                              19478 non-null  int64
11  JobRole                               19478 non-null  object
12  JobSatisfaction                       19478 non-null  int64
13  MaritalStatus                         19478 non-null  object
14  MonthlyIncome                        19478 non-null  int64
15  MonthlyRate                          19478 non-null  int64
16  NumCompaniesWorked                   19478 non-null  int64
17  OverTime                             19478 non-null  object
18  PercentSalaryHike                    19478 non-null  int64
19  RelationshipSatisfaction              19478 non-null  int64
20  StandardHours                        19478 non-null  int64
21  StockOptionLevel                     19478 non-null  int64
22  TotalWorkingYears                    19478 non-null  int64
23  TrainingTimesLastYear                19478 non-null  int64
24  WorkLifeBalance                      19478 non-null  int64
25  YearsAtCompany                       19478 non-null  int64
26  YearsInCurrentRole                   19478 non-null  int64
27  YearsSinceLastPromotion              19478 non-null  int64
28  YearsWithCurrManager                 19478 non-null  int64
29  Employee Source                       19478 non-null  object
dtypes: int64(23), object(7)
memory usage: 4.5+ MB
```

in this dataset we have 19478 rows and 30 columns in which 23 are int datatype and 7 are object with no null values present

```
df.describe()
```

	Age	DailyRate	DistanceFromHome	Education	EnvironmentSatisfac
count	19478.000000	19478.000000	19478.000000	19478.000000	19478.00
mean	37.524489	812.367697	8.945836	2.924274	2.74
std	8.860420	402.778087	8.006485	1.026008	1.08
min	18.000000	102.000000	1.000000	1.000000	1.00
25%	31.000000	477.000000	2.000000	2.000000	2.00
50%	36.000000	813.000000	7.000000	3.000000	3.00
75%	43.000000	1176.000000	13.000000	4.000000	4.00
max	60.000000	1499.000000	29.000000	5.000000	4.00

8 rows × 23 columns

✓ Now lets find null values

```
df.isnull().sum()/len(df)*100
```

```
Age                0.0
BusinessTravel     0.0
DailyRate          0.0
Department         0.0
DistanceFromHome   0.0
Education           0.0
EnvironmentSatisfaction  0.0
Gender             0.0
HourlyRate         0.0
JobInvolvement     0.0
JobLevel           0.0
JobRole            0.0
JobSatisfaction    0.0
MaritalStatus      0.0
MonthlyIncome      0.0
MonthlyRate        0.0
NumCompaniesWorked 0.0
OverTime           0.0
PercentSalaryHike  0.0
RelationshipSatisfaction  0.0
StandardHours      0.0
StockOptionLevel   0.0
TotalWorkingYears  0.0
TrainingTimesLastYear  0.0
WorkLifeBalance    0.0
YearsAtCompany     0.0
YearsInCurrentRole 0.0
YearsSinceLastPromotion 0.0
YearsWithCurrManager 0.0
Employee Source    0.0
dtype: float64
```

✓ There is no null values present in this dataset

```
for i in df:
    print(i)
    print(df[i].unique())
```

```
[3 2 4 0 1 0]
WorkLifeBalance
[3 2 4 1]
YearsAtCompany
[10 14 6 0 8 3 5 2 1 37 7 4 20 17 9 15 16 11 12 13 22 18 25 21
 27 40 33 24 19 36 29 31 32 26 30 34 23]
YearsInCurrentRole
[ 7 8 5 0 2 3 10 4 15 1 9 13 11 6 12 16 14 18 17]
YearsSinceLastPromotion
[ 1 6 0 3 2 4 7 5 8 12 13 11 10 9 15 14]
YearsWithCurrManager
[ 7 9 3 8 0 2 4 6 12 5 15 1 10 11 17 13 16 14]
Employee Source
['Seek' 'Indeed' 'Referral' 'Company Website' 'Adzuna' 'GlassDoor' 'Jora'
 'LinkedIn' 'Recruit.net']
```

▾ Separating cat and num columns

```
cat_col=df.select_dtypes(object)
cat_col
```

	BusinessTravel	Department	Gender	JobRole	MaritalStatus	OverTime	Er
0	Travel_Frequently	Research & Development	Male	Research Scientist	Married	No	
1	Non-Travel	Human Resources	Male	Research Scientist	Married	No	
2	Non-Travel	Human Resources	Male	Research Scientist	Married	No	
3	Travel_Frequently	Research & Development	Male	Research Scientist	Married	No	
4	Travel_Frequently	Research & Development	Male	Research Scientist	Married	No	
...	
19473	Travel_Rarely	Research & Development	Female	Laboratory Technician	Single	No	
19474	Travel_Rarely	Research & Development	Female	Research Scientist	Single	No	

```
num_col=df.select_dtypes([int,float])
num_col
```

	Age	DailyRate	DistanceFromHome	Education	EnvironmentSatisfaction	HourlyRate
0	49	279	8	1	3	
1	59	1420	2	4	1	
2	59	1420	2	4	1	
3	49	279	8	1	3	
4	49	279	8	1	3	
...	
19473	47	465	1	3	4	
19474	38	371	2	3	4	
19475	34	629	27	2	4	
19476	55	177	8	1	4	
19477	27	1134	16	4	4	

19478 rows × 23 columns

```
for i in cat_col:
    print(i)
    print(cat_col[i].unique())

BusinessTravel
['Travel_Frequently' 'Non-Travel' 'Travel_Rarely']
Department
['Research & Development' 'Human Resources' 'Sales']
Gender
['Male' 'Female']
JobRole
```

```

['Research Scientist' 'Manufacturing Director' 'Laboratory Technician'
 'Sales Representative' 'Sales Executive' 'Manager' 'Human Resources'
 'Healthcare Representative' 'Research Director']
MaritalStatus
['Married' 'Single' 'Divorced']
OverTime
['No' 'Yes']
Employee Source
['Seek' 'Indeed' 'Referral' 'Company Website' 'Adzuna' 'GlassDoor' 'Jora'
 'LinkedIn' 'Recruit.net']

for i in num_col:
    print(i)
    print(num_col[i].unique())

1448 601 1221 383 1109 264 918 788 1313 1186 1464 196 796 723
415 337 937 1492 801 704 301 1120 469 1262 1308 984 174 718
367 1384 902 669 1457 1421 150 179 363 107 1465 1098 969 1320
1429 603 968 879 640 266 412 1138 325 634 1253 1202 256 1405
999 285 404 683 1462 949 652 332 560 359 866 1326 748 990
1193 271 333 1440 674 441 1342 898 350 992 1288 1108 479 1059
457 241 1015 1387 1470 365 486 1037 392 567 148 786 370 146
611 897 1054 181 734 1128 1180 431 572 352 1172 1079 1394 1239
911 1162 234 468 613 1023 628 590 953 355 835 219 1096 1444
1382 1378 1266 529]
DistanceFromHome
[ 8  2 26  3  9 10 13 23 18 24  1  5  7 27  6 25 20 16 15 19 21  4 11 29
 22 28 14 12 17]
Education
[1 4 3 2 5]
EnvironmentSatisfaction
[3 1 2 4]
HourlyRate
[ 61  37  95  87  56  79  62  40  33  57  42  81  67  90  44  66  53  31
  94  55  93  84  49  38  32  52  69  86  70  30  50  51  88  80  96  65
  78  45  46  41  82  99  58  39  48  63  72  83  97  75  73  98  36  47
  71  77  43  59  76  60  54 100  35  64  92  91  34  89  68  85  74]
JobInvolvement
[2 3 1 4]
JobLevel
[2 1 5 3 4]
JobSatisfaction
[2 3 1 4]
MonthlyIncome
[5130 6499 4810 ... 9991 5390 4404]
MonthlyRate
[24907 22656 26314 ... 5174 13243 10228]
NumCompaniesWorked
[1 2 3 6 4 9 0 7 8 5]
PercentSalaryHike
[23 13 14 15 11 16 12 17 21 20 22 18 19 24 25]
RelationshipSatisfaction
[4 3 2 1]
StandardHours
[80]
StockOptionLevel
[1 0 2 3]
TotalWorkingYears
[10 16  6 19  7  8  5  2 38  3 20 17  1 12 22 21 13  9 14 11  4 30 23 28
 15 32 24 31 26 37  0 40 29 18 25 34 36 35 33 27]
TrainingTimesLastYear
[3 5 2 4 0 1 6]
WorkLifeBalance
[3 2 4 1]
YearsAtCompany
[10 14  6  0  8  3  5  2  1 37  7  4 20 17  9 15 16 11 12 13 22 18 25 21
 27 40 33 24 19 36 29 31 32 26 30 34 23]
YearsInCurrentRole
[ 7  8  5  0  2  3 10  4 15  1  9 13 11  6 12 16 14 18 17]
YearsSinceLastPromotion
[ 1  6  0  3  2  4  7  5  8 12 13 11 10  9 15 14]
YearsWithCurrManager
[ 7  9  3  8  0  2  4  6 12  5 15  1 10 11 17 13 16 14]

```

✓ Encoding

```

from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()

```

```
for i in cat_col:
    cat_col[i]=le.fit_transform(cat_col[i])
    print(cat_col[i].unique())
    print(le.classes_)

[1 0 2]
['Non-Travel' 'Travel_Frequently' 'Travel_Rarely']
[1 0 2]
['Human Resources' 'Research & Development' 'Sales']
[1 0]
['Female' 'Male']
[6 4 2 8 7 3 1 0 5]
['Healthcare Representative' 'Human Resources' 'Laboratory Technician'
 'Manager' 'Manufacturing Director' 'Research Director'
 'Research Scientist' 'Sales Executive' 'Sales Representative']
[1 2 0]
['Divorced' 'Married' 'Single']
[0 1]
['No' 'Yes']
[8 3 7 1 0 2 4 5 6]
['Adzuna' 'Company Website' 'GlassDoor' 'Indeed' 'Jora' 'LinkedIn'
 'Recruit.net' 'Referral' 'Seek']
```

```
cat_col.head()
```

	BusinessTravel	Department	Gender	JobRole	MaritalStatus	OverTime	Employee Source
0	1	1	1	6	1	0	8
1	0	0	1	6	1	0	8
2	0	0	1	6	1	0	8
3	1	1	1	6	1	0	8
4	1	1	1	6	1	0	8

```
new_df=pd.concat([cat_col,num_col],axis=1)
```

```
new_df
```

	BusinessTravel	Department	Gender	JobRole	MaritalStatus	OverTime	Employee Source
0	1	1	1	6	1	0	8
1	0	0	1	6	1	0	8
2	0	0	1	6	1	0	8
3	1	1	1	6	1	0	8
4	1	1	1	6	1	0	8
...
19473	2	1	0	2	2	0	4
19474	2	1	1	6	2	1	1
19475	2	1	0	6	2	0	1
19476	0	1	0	4	1	1	6
19477	2	1	1	1	1	1	1

19478 rows × 30 columns

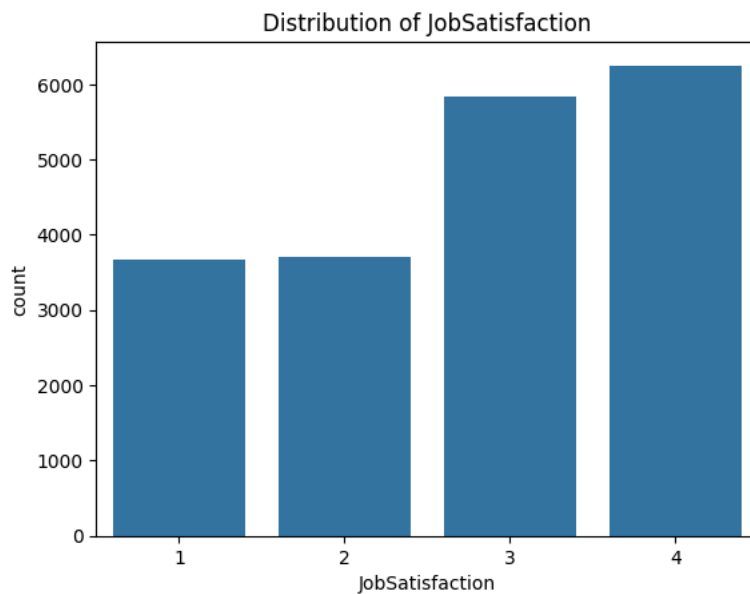
```
new_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19478 entries, 0 to 19477
Data columns (total 30 columns):
#   Column                Non-Null Count  Dtype
---  -
0   BusinessTravel        19478 non-null  int64
1   Department            19478 non-null  int64
2   Gender                19478 non-null  int64
3   JobRole               19478 non-null  int64
4   MaritalStatus         19478 non-null  int64
5   OverTime              19478 non-null  int64
6   Employee Source       19478 non-null  int64
7   Age                  19478 non-null  int64
8   DailyRate             19478 non-null  int64
9   DistanceFromHome      19478 non-null  int64
```

```
10 Education          19478 non-null int64
11 EnvironmentSatisfaction  19478 non-null int64
12 HourlyRate          19478 non-null int64
13 JobInvolvement      19478 non-null int64
14 JobLevel            19478 non-null int64
15 JobSatisfaction      19478 non-null int64
16 MonthlyIncome        19478 non-null int64
17 MonthlyRate          19478 non-null int64
18 NumCompaniesWorked   19478 non-null int64
19 PercentSalaryHike    19478 non-null int64
20 RelationshipSatisfaction  19478 non-null int64
21 StandardHours        19478 non-null int64
22 StockOptionLevel     19478 non-null int64
23 TotalWorkingYears    19478 non-null int64
24 TrainingTimesLastYear  19478 non-null int64
25 WorkLifeBalance      19478 non-null int64
26 YearsAtCompany       19478 non-null int64
27 YearsInCurrentRole    19478 non-null int64
28 YearsSinceLastPromotion  19478 non-null int64
29 YearsWithCurrManager  19478 non-null int64
dtypes: int64(30)
memory usage: 4.5 MB
```

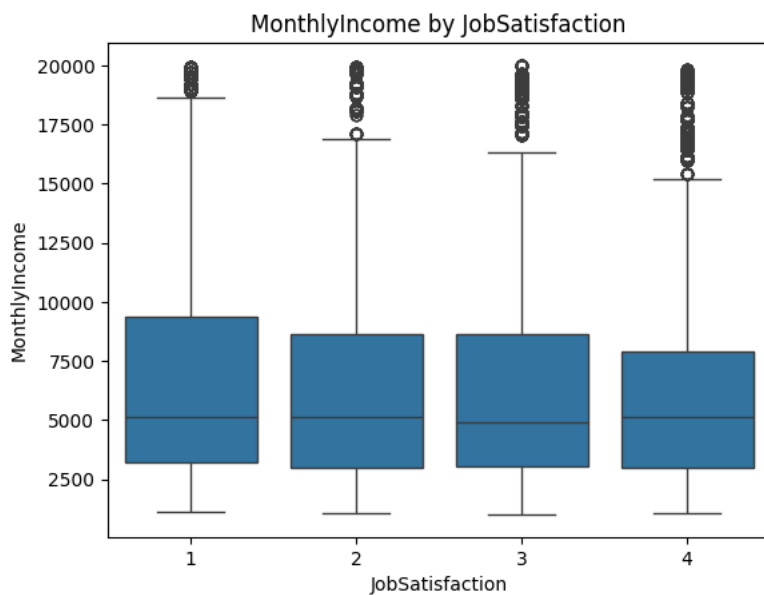
EDA

```
sns.countplot(data=df, x='JobSatisfaction')
plt.title('Distribution of JobSatisfaction')
plt.show()
```




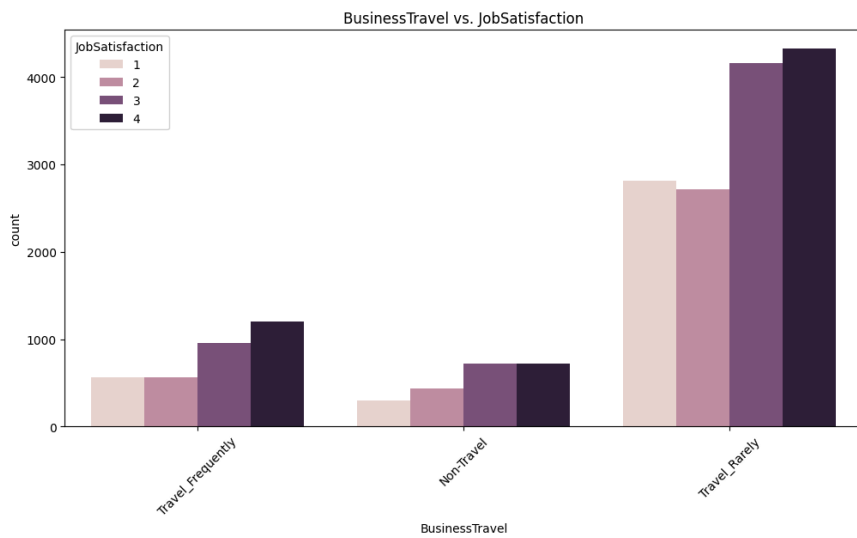
```
sns.histplot(data=df, x='Age', hue='JobSatisfaction', kde=True)
plt.title('Age Distribution by JobSatisfaction')
plt.show()
```

```
sns.boxplot(data=df, x='JobSatisfaction', y='MonthlyIncome')
plt.title('MonthlyIncome by JobSatisfaction')
plt.show()
```

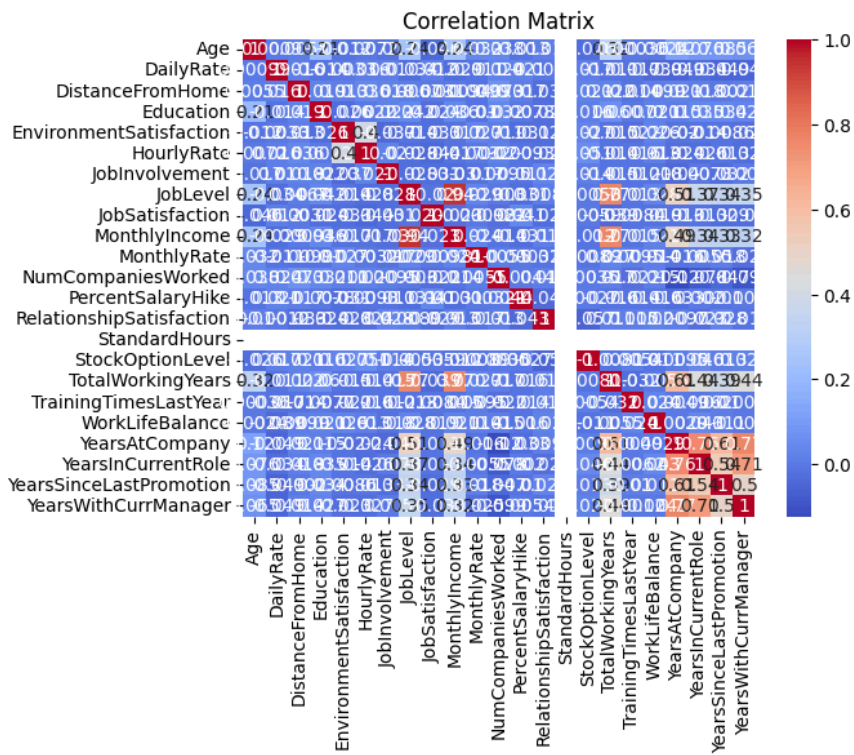


```
plt.figure(figsize=(12, 6))
sns.countplot(data=df, x='BusinessTravel', hue='JobSatisfaction')
plt.title('BusinessTravel vs. JobSatisfaction')
plt.xticks(rotation=45)
plt.show()
```

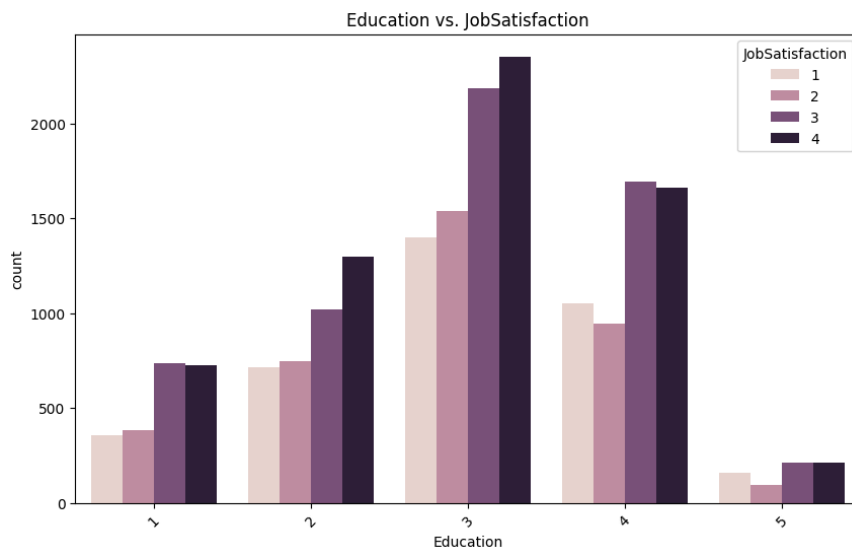
```
plt.figure(figsize=(10, 6))
sns.countplot(data=df, x='JobRole', hue='JobSatisfaction')
plt.title('JobRole vs. JobSatisfaction')
plt.xticks(rotation=45)
plt.show()
```

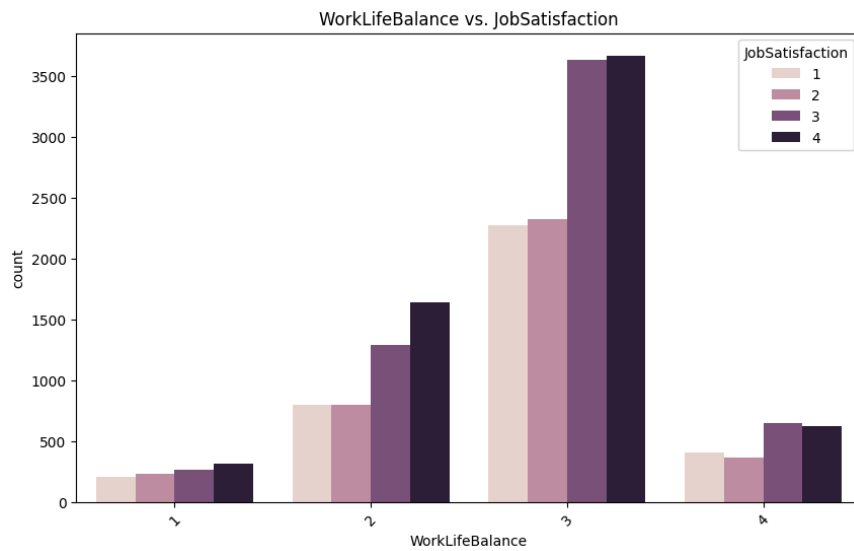
```
correlation_matrix = df.corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



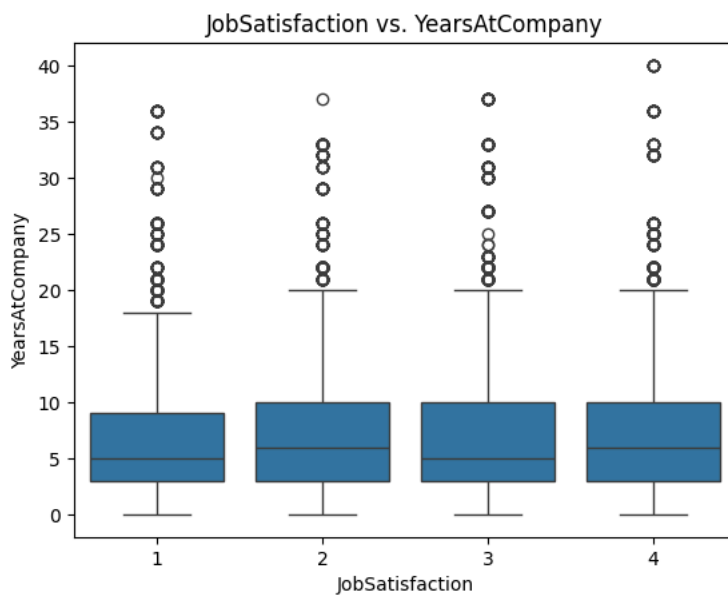
```
plt.figure(figsize=(10, 6))
sns.countplot(data=df, x='Education', hue='JobSatisfaction')
plt.title('Education vs. JobSatisfaction')
plt.xticks(rotation=45)
plt.show()
```



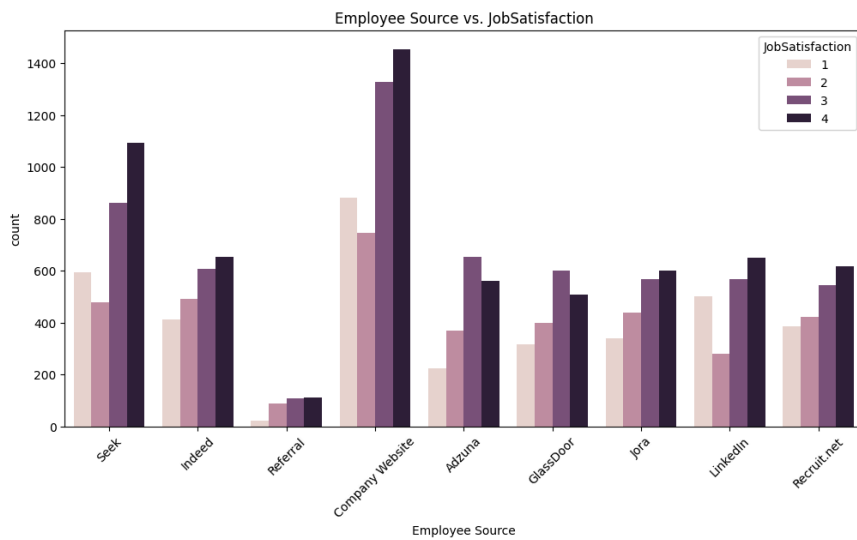
```
plt.figure(figsize=(10, 6))
sns.countplot(data=df, x='WorkLifeBalance', hue='JobSatisfaction')
plt.title('WorkLifeBalance vs. JobSatisfaction')
plt.xticks(rotation=45)
plt.show()
```



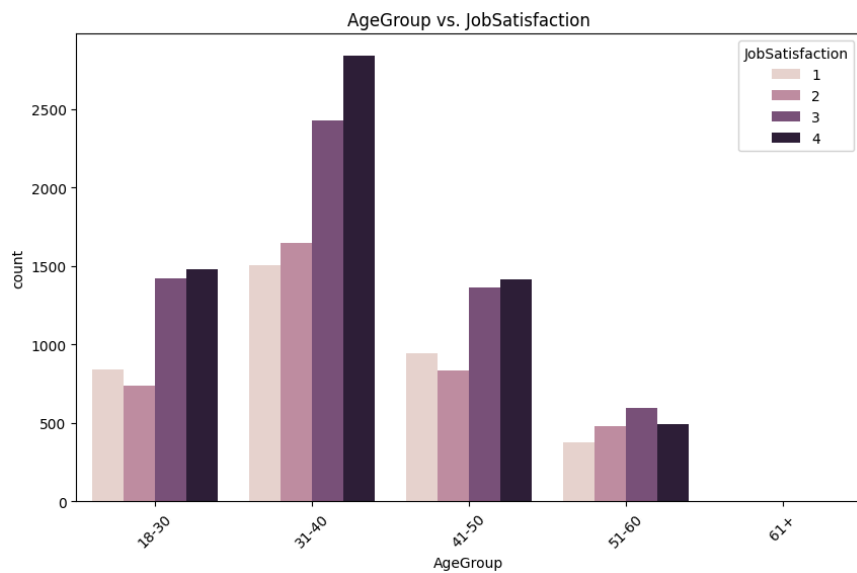
```
sns.boxplot(data=df, x='JobSatisfaction', y='YearsAtCompany')
plt.title('JobSatisfaction vs. YearsAtCompany')
plt.show()
```



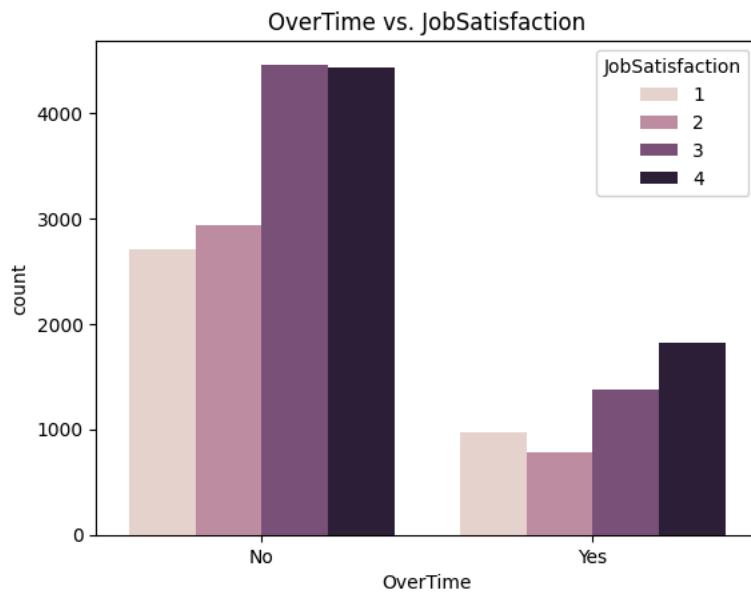
```
plt.figure(figsize=(12, 6))
sns.countplot(data=df, x='Employee Source', hue='JobSatisfaction')
plt.title('Employee Source vs. JobSatisfaction')
plt.xticks(rotation=45)
plt.show()
```



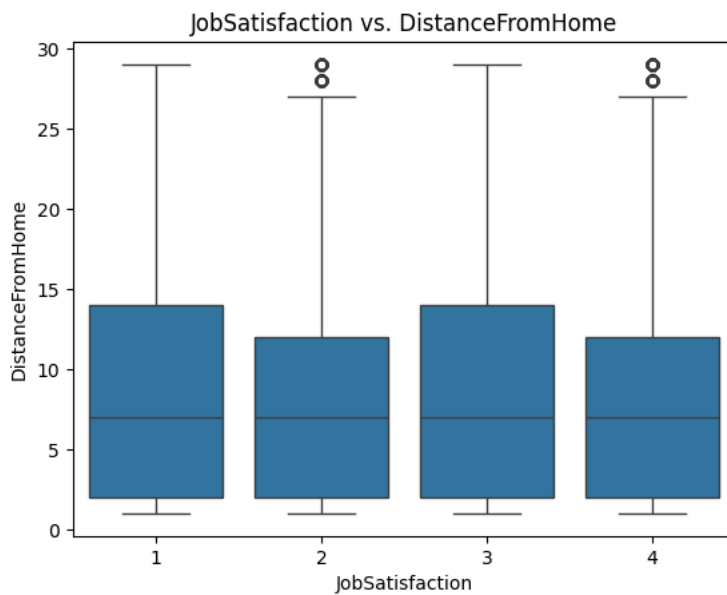
```
df['AgeGroup'] = pd.cut(df['Age'], bins=[18, 30, 40, 50, 60, np.inf], labels=['18-30', '31-40', '41-50', '51-60', '61+'])
plt.figure(figsize=(10, 6))
sns.countplot(data=df, x='AgeGroup', hue='JobSatisfaction')
plt.title('AgeGroup vs. JobSatisfaction')
plt.xticks(rotation=45)
plt.show()
```



```
sns.countplot(data=df, x='OverTime', hue='JobSatisfaction')
plt.title('OverTime vs. JobSatisfaction')
plt.show()
```




```
sns.boxplot(data=df, x='JobSatisfaction', y='DistanceFromHome')
plt.title('JobSatisfaction vs. DistanceFromHome')
plt.show()
```



✓ Feature selection

```
plt.figure(figsize=(20,10))
sns.heatmap(new_df.corr(),annot=True,cmap='coolwarm',fmt=".3f")
```

<Axes: >




```
corr=new_df.corr()['JobSatisfaction'].reset_index()
corr.sort_values('JobSatisfaction',ascending=False)
```

	index	JobSatisfaction
15	JobSatisfaction	1.000000
4	MaritalStatus	0.048312
19	PercentSalaryHike	0.040541
5	OverTime	0.034882
2	Gender	0.028142
3	JobRole	0.020238
1	Department	0.012380
8	DailyRate	0.011897
6	Employee Source	-0.000816
17	MonthlyRate	-0.000979
9	DistanceFromHome	-0.003112
22	StockOptionLevel	-0.005881
24	TrainingTimesLastYear	-0.008429
27	YearsInCurrentRole	-0.012658
26	YearsAtCompany	-0.013406
25	WorkLifeBalance	-0.018608
14	JobLevel	-0.020063
18	NumCompaniesWorked	-0.021878
16	MonthlyIncome	-0.023214
10	Education	-0.024319
20	RelationshipSatisfaction	-0.028806
28	YearsSinceLastPromotion	-0.029380
29	YearsWithCurrManager	-0.030053
13	JobInvolvement	-0.031465
11	EnvironmentSatisfaction	-0.032697
23	TotalWorkingYears	-0.038994
12	HourlyRate	-0.043590
7	Age	-0.046104
0	BusinessTravel	-0.055352
21	StandardHours	NaN

```

corelation = pd.DataFrame(new_df.corr())
corelation = pd.DataFrame(corelation['JobSatisfaction'])
corelation=corelation.sort_values('JobSatisfaction',ascending=False)
indices_to_remove = ['JobSatisfaction']
corelation = corelation.drop(indices_to_remove)
plt.figure(figsize=(8,8))
sns.barplot(x=corelation['JobSatisfaction'],y=corelation.index)

```



```
new_df.columns

Index(['BusinessTravel', 'Department', 'Gender', 'JobRole', 'MaritalStatus',
      'OverTime', 'Employee Source', 'Age', 'DailyRate', 'DistanceFromHome',
      'Education', 'EnvironmentSatisfaction', 'HourlyRate', 'JobInvolvement',
      'JobLevel', 'JobSatisfaction', 'MonthlyIncome', 'MonthlyRate',
      'NumCompaniesWorked', 'PercentSalaryHike', 'RelationshipSatisfaction',
      'StandardHours', 'StockOptionLevel', 'TotalWorkingYears',
      'TrainingTimesLastYear', 'WorkLifeBalance', 'YearsAtCompany',
      'YearsInCurrentRole', 'YearsSinceLastPromotion',
      'YearsWithCurrManager'],
      dtype='object')

selected_columns = ['MaritalStatus', 'PercentSalaryHike', 'OverTime', 'Gender', 'JobRole', 'EnvironmentSatisf
sat= new_df[selected_columns]

sat
```

	MaritalStatus	PercentSalaryHike	OverTime	Gender	JobRole	EnvironmentSatisf
0	1	23	0	1	6	
1	1	23	0	1	6	
2	1	23	0	1	6	
3	1	23	0	1	6	
4	1	23	0	1	6	
...
19473	2	13	0	0	2	
19474	2	11	1	1	6	
19475	2	15	0	0	6	
19476	1	14	1	0	4	
19477	1	11	1	1	1	

19478 rows × 12 columns

✓ Feature Selection

```
X=sat.drop('JobSatisfaction',axis=1)
X
```

	MaritalStatus	PercentSalaryHike	OverTime	Gender	JobRole	EnvironmentSatisf
0	1	23	0	1	6	
1	1	23	0	1	6	
2	1	23	0	1	6	
3	1	23	0	1	6	
4	1	23	0	1	6	
...
19473	2	13	0	0	2	
19474	2	11	1	1	6	
19475	2	15	0	0	6	
19476	1	14	1	0	4	
19477	1	11	1	1	1	


19478 rows x 11 columns

```
y=sat['JobSatisfaction']
y
```

```
0      2
1      2
2      2
3      2
4      2
..
19473  2
19474  4
19475  2
19476  1
19477  2
Name: JobSatisfaction, Length: 19478, dtype: int64
```

```
sns.countplot(df['JobSatisfaction'])
```

<Axes: ylabel='count'>



✓ Train test split

```
from sklearn.model_selection import train_test_split
```

```
Xtrain,Xtest,ytrain,ytest=train_test_split(X,y,stratify=y,test_size=0.2,random_state=42)
```

```
Xtrain.shape , Xtest.shape

((15582, 11), (3896, 11))

ytrain.shape , ytest.shape

((15582,), (3896,))
```

✓ Scaling

```
from sklearn.preprocessing import StandardScaler
se=StandardScaler()
```

```
Xtrain=se.fit_transform(Xtrain)
Xtest=se.fit_transform(Xtest)
```

✓ Training models

```
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import AdaBoostClassifier

from sklearn.metrics import accuracy_score,classification_report
from sklearn.metrics import roc_auc_score,roc_curve
```

```
knn=KNeighborsClassifier(n_neighbors=3)
lr=LogisticRegression()
dt=DecisionTreeClassifier()
ra=RandomForestClassifier()
ad=AdaBoostClassifier()
svm=SVC(probability=True)
gau=GaussianNB()
bag=BaggingClassifier()
Gr=GradientBoostingClassifier()
```

```
Training_score= []
Testing_score= []
def model_building(model):
    model.fit(Xtrain, ytrain)
    ytrain_pred= model.predict(Xtrain)
    ytest_pred= model.predict(Xtest)
    a= accuracy_score(ytrain, ytrain_pred)
    b= accuracy_score(ytest, ytest_pred)
    Training_score.append(a)
    Testing_score.append(b)
    print(model)
    print("Train Data\n", accuracy_score(ytrain,ytrain_pred))
    print("Test Data\n", accuracy_score(ytest,ytest_pred))
```

```
model_building(knn)

KNeighborsClassifier(n_neighbors=3)
Train Data
0.9896675651392632
Test Data
0.981776180698152
```

```
model_building(lr)

LogisticRegression()
Train Data
0.34539853677319987
Test Data
0.3408624229979466
```

```
model_building(dt)

    DecisionTreeClassifier()
    Train Data
    1.0
    Test Data
    0.8986139630390144

model_building(ra)

    RandomForestClassifier()
    Train Data
    1.0
    Test Data
    0.9961498973305954

model_building(ad)

    AdaBoostClassifier()
    Train Data
    0.3762674881273264
    Test Data
    0.36113963039014374

model_building(svm)

    SVC(probability=True)
    Train Data
    0.6419586702605571
    Test Data
    0.6054928131416838

model_building(gau)

    GaussianNB()
    Train Data
    0.340007701193685
    Test Data
    0.32931211498973306

model_building(bag)

    BaggingClassifier()
    Train Data
    0.9998074701578745
    Test Data
    0.9840862422997947

model_building(Gr)

    GradientBoostingClassifier()
    Train Data
    0.6336798870491593
    Test Data
    0.48716632443531827

Models= ["k-Nearest Neighbors","Logistic Regression" ,"Decision Tree Classifier", "Random forest Classifier" ,
        "Ada-Boosting Classifier","svm","GaussianNB","Bagging Classifier", "Gradient- Bossting Classifier"]

new_df1 = pd.DataFrame({"Algorithms":Models,
                        "Training Score":Training_score,
                        "Testing Score":Testing_score,})

new_df1
```

	Algorithms	Training Score	Testing Score
5	svm	0.641959	0.605493

Hypertunning

Random forest

```

5                                svm                0.641959                0.605493

from sklearn.model_selection import RandomizedSearchCV

# Random Forest Classifier
ra=RandomForestClassifier()

# Random Forest Parameters
random_forest_params = {
    'n_estimators': [25,50,75,100],
    'max_depth': [2, 3, 5, 10, 20],
    'min_samples_leaf': [5, 10, 20, 50, 100],
    'min_samples_split': [2, 5, 10],
    'criterion': ["gini", "entropy"],
    'max_features': ['auto', 'sqrt'],
    'bootstrap': [True, False],
    'class_weight' : ["balanced", "balanced_subsample"]

}

ra_reg=RandomizedSearchCV(ra,param_distributions=random_forest_params,random_state=42,scoring='accuracy',cv=5,n_jobs=-1)

model_building(ra_reg)

RandomizedSearchCV(cv=5, estimator=RandomForestClassifier(), n_jobs=-1,
                    param_distributions={'bootstrap': [True, False],
                                         'class_weight': ['balanced',
                                                         'balanced_subsample'],
                                         'criterion': ['gini', 'entropy'],
                                         'max_depth': [2, 3, 5, 10, 20],
                                         'max_features': ['auto', 'sqrt'],
                                         'min_samples_leaf': [5, 10, 20, 50,
                                                             100],
                                         'min_samples_split': [2, 5, 10],
                                         'n_estimators': [25, 50, 75, 100]},
                    random_state=42, scoring='accuracy')

Train Data
0.835515338210756
Test Data
0.7517967145790554

```

gaussianNB

```

from sklearn.model_selection import GridSearchCV
from sklearn.naive_bayes import GaussianNB

```